

PROGRAMMING LANGUAGES FOR DISTRIBUTED SYSTEMS



From Languages
to Services

Prof. P. Felber
Pascal.Felber@unine.ch
<http://iiun.unine.ch/>



Our Group

- “Complex Systems” research group
 - ~10 people (~7 PhD students, ~3 senior)
- Research topics
 - Dependability (fault tolerance...)
 - Distributed systems (Cloud, IoT...)
 - Concurrency (multi-core, transactional memory...)
 - Big Data (processing, storage, NVM...)
 - Energy (power-efficient computing...)
 - Security (TEEs, storage, blockchain...)
- Funding: mainly EU and Swiss projects

Our Experience with PL4DS

SPLAY

- SPLAY: “*distributed applications made simple*” [NSDI’09]
 - Goal: quick prototyping → real deployment/evaluation
 - Easy to read and develop (\approx pseudo-code), lightweight
 - Used extensively in teaching since 2011
- Principle
 - Sandboxed environment deployed on remote nodes
 - Based on Lua (event-driven + coroutines)
 - Collection of libraries (networking, logging, serialization, storage, crypto...)
 - Reproducible experiments through controller (trace replay, churn generator with DSL)
 - Easily extensible, open-source [<http://splay-project.org>]

From Pseudo-code to Real Code...

```

// join a Chord ring containing node n'.
n.join(n')
  predecessor = nil;
  successor = n'.find_successor(n);

// called periodically: verifies n's immediate
// successor, and tells the successor about n.
n.stabilize()
  x = successor.predecessor;
  if (x ∈ (n, successor))
    successor = x;
    successor.notify(n);

// n' thinks it might be our predecessor.
n.notify(n')
  if (predecessor is nil or n' ∈ (predecessor, n))
    predecessor = n';

// called periodically: refreshes finger table entries.
// next stores the index of the next finger to fix.
n.fix_fingers()
  next = next + 1;
  if (next > m)
    next = 1;
  finger[next] = find_successor(n + 2ext-1);

// called periodically: checks whether predecessor has failed.
n.check_predecessor()
  if (predecessor has failed)
    predecessor = nil;

// ask node n to find the successor of id
n.find_successor(id)
  if (id ∈ (n, successor])
    return successor;
  else
    n' = closest_preceding_node(id);
    return n'.find_successor(id);

// search the local table for the highest predecessor of id
n.closest_preceding_node(id)
  for i = m downto 1
    if (finger[i] ∈ (n, id))
      return finger[i];
  return n;

```



```

1  function join(n0) -- n0: some node in the ring
2    predecessor = nil
3    finger[1] = rpc(n0, {'find_successor', {n.id}})
4  end

5  function stabilize() -- periodically verify n's successor
6    local x = rpc(finger[1], {'predecessor'})
7    if between(x.id, n.id, finger[1].id) then
8      finger[1] = x -- new successor
9    end
10   rpc(finger[1], {'notify', {n}})
11 end

12 function notify(n0) -- n0 thinks it might be our predecessor
13   if not predecessor or between(n0.id, predecessor.id, n.id) then
14     predecessor = n0 -- new predecessor
15   end
16 end

17 function fix_fingers() -- refresh fingers
18   next = (next % m) + 1 -- 1 ≤ next ≤ m
19   finger[next] = find_successor((n.id + 2^(next - 1)) % 2^m)
20 end

21 function check_predecessor() -- checks if predecessor has failed
22   if predecessor and not rpc(predecessor) then
23     predecessor = nil
24   end
25 end

25 function find_successor(id) -- ask node to find id's successor
26   if between(id, n.id, (finger[1].id + 1) % 2^m) then
27     return finger[1]
28   end
29   local n0 = closest_preceding_node(id)
30   return rpc(n0, {'find_successor', {id}})
31 end

32 function closest_preceding_node(id) -- finger preceding id
33   for i = m, 1, -1 do
34     if finger[i] and between(finger[i].id, n.id, id) then
35       return finger[i]
36     end
37   end
38   return n
39 end

```

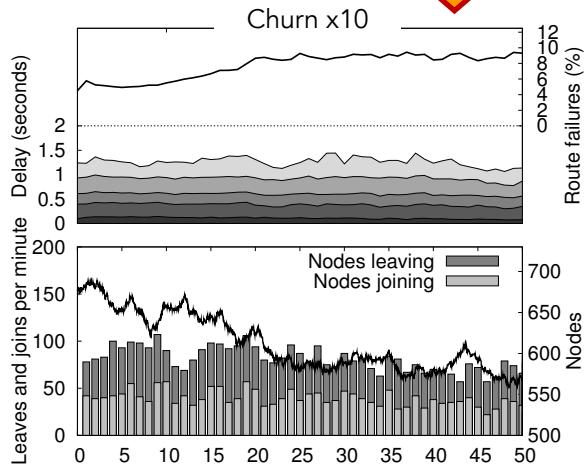
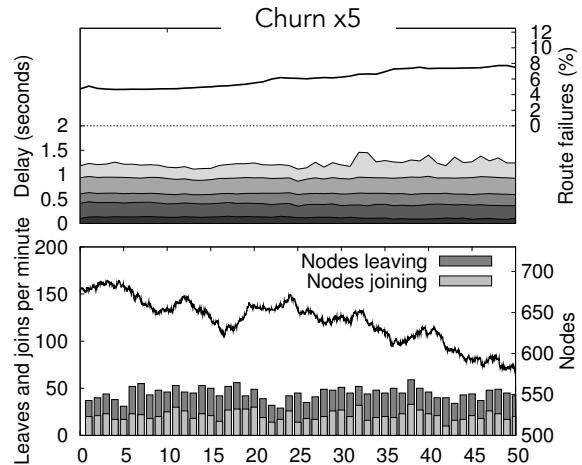
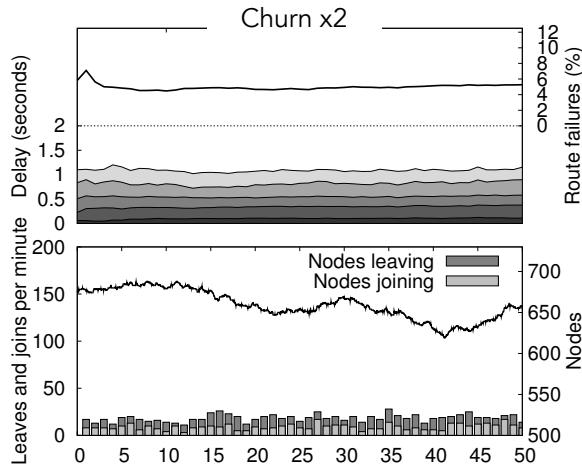
...to Deployment and Evaluation



VALUES	
NUMBER OF SPLAYS	<input type="text" value="25"/>
TRACE (NUMBER OF LINES => NUMBER OF SPLAYS)	<input type="text"/> Parcourir...
MAX EXECUTION TIME (SECONDS)	<input type="text" value="10000"/>
SPLAYD'S VERSION	<input type="text" value="0.881"/>
MAX LOAD	<input type="text" value="5"/>
MIN UPTIME (SECONDS)	<input type="text" value="0"/>
COUNTRY/CONTINENT CODE (2 LETTERS)	<input type="text"/>
GEOLOCALIZATION (LATITUDE / LONGITUDE / DISTANCE (KM))	<input type="text" value="49.610709936"/> / <input type="text" value="5.44921875"/> / <input type="text" value="1700"/>
Click on the map to set the center or center to your position (and remember to set a distance).	
Plan	Satellite
Mixte	




Pastry on PlanetLab with Overnet traces



Outcome and Lessons Learned

- Several extensions
 - SPLAYNET: User-Space Topology Emulation [Middleware'13]
 - Teaching material, DSLs
- Pros
 - Compact, easy development, fast deployment, reproducible experiments, improved “time to paper”
 - Good for research and testing algorithms “in the wild”
- Cons
 - Corner cases hard to handle (simplicity vs. flexibility)
 - **Not good enough for production systems and complete applications** (but great for prototyping)
 - Limited security!

Do we Need New Languages?

- Many languages are domain-specific
 - Stream processing, dataflow, protocols, HPC, etc.
- Distributed systems are complex
 - Composition of several remote components

Single-language

vs.

composition of services

(each one developed with the most appropriate language)

Do we Need New Languages?

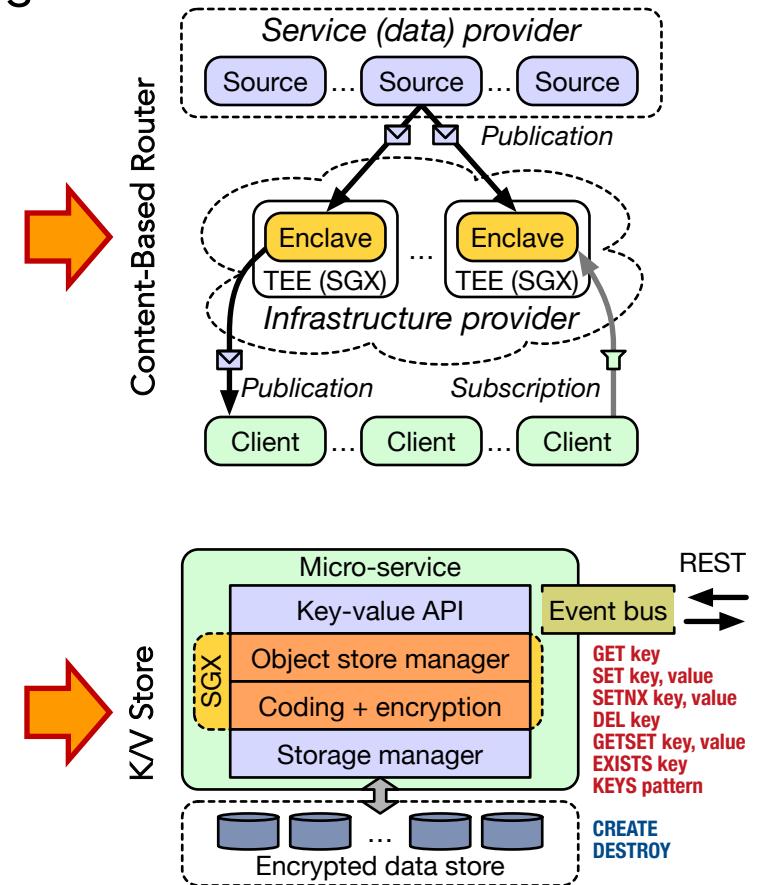
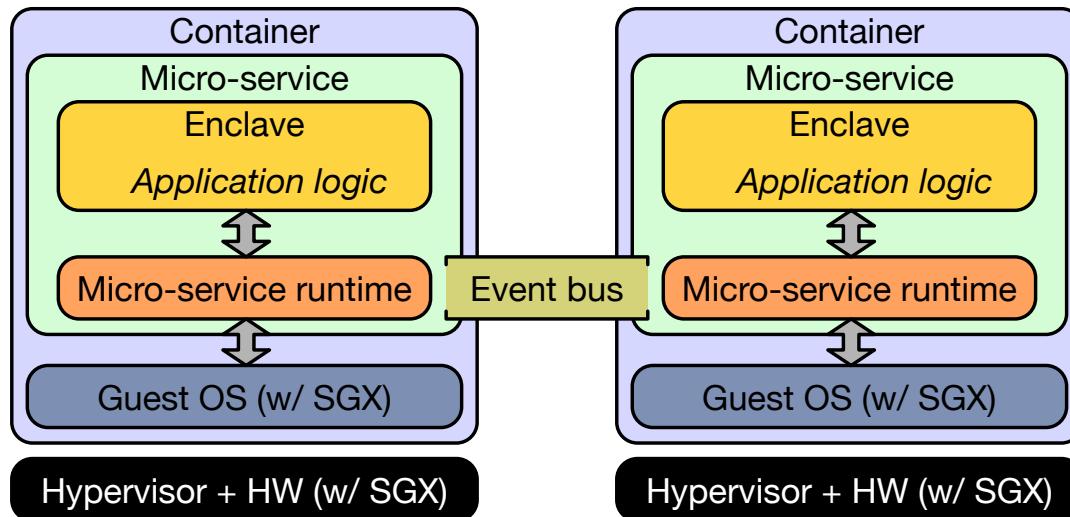
1. Microservices are great for distributed systems
 - Partition the application according to its components
 - Partition the developer teams according to their skills
 - Can be easily *containerized*
2. We need to support multiple languages with a simple way to interface the microservices
 - REST APIs are great for that: standard, interoperable
3. We need support for “composition”
(orchestration, coordination, scheduling...)
 - Many tools available (e.g., Kubernetes, Swarm)

The Quest for Security

- Production systems must be protected
 - Mission-critical, vulnerable to hackers
 - Manage sensitive data
- Distributed systems are exposed
- Remote code must be secure
 - Protect the environment from the application
 - Protect the application from the environment
- PL4DS must support development of dependable (\approx trustworthy) software
- Execution environment must be secure (TEE)

Our Experience with Secure μ S

- SECURECLOUD: “Secure Distributed Big Data Application with Micro-Services” [H2020 EU-Brazil]
 - Micro-services within containers in TEEs in the cloud



Outcome and Lessons Learned

- Developed many μ-services (building blocks)
 - Communication, storage, processing (map-reduce)
 - Services developed in several languages (C/C++, Go, Rust, Java, Python, Lua...)
 - Run securely in containers in public clouds within TEEs
- Pros
 - Best language for each problem
 - Communication/distribution using standard tools
 - Easy composition
- Cons
 - No single language/development environment

Thanks!

DISCUSSION?