






Hidehiko Masuhara
Tokyo Tech
(with Matthias Springer)

A Parallel Object-Oriented Programming Language for GPGPU



Concurrency model of GPGPU



- At the low levels
 - SIMD (synchronized) execution
 - Distributed memory
- At the higher levels
 - No (or only coarse) synchronization
 - Globally shared memory
- State-of-the-art of programming style
 - Static data allocation + data-parallel

Motivation: wanted a quick prototyping language for GPGPU



interested in GPGPU algorithms
but CUDA/C is too naive...



a scripting language for GPGPU!?



... +




Ikra: data-parallel extension to Ruby




- Concurrency model:
parallel arrays with map & reduce
- OOP support
 - Optimized memory layout for GPU
 - Parallel arrays (or sets) of objects
 - Dynamic object allocation ←new!

[ECOOP'19, ISMM'19]


Parallel OOP Applications (demo)




(a) barnes-hut [4]:
Parallel Tree Constr.




(b) collisions:
Particle System




(c) game-of-life:
Cellular Automaton



(d) nbody:
Particle System




(e) structure [14]:
Finite Elem. Method



(f) sugarscape [8]:
Agent-based Sim.



(g) traffic [17]:
Nagel-Schr. Model

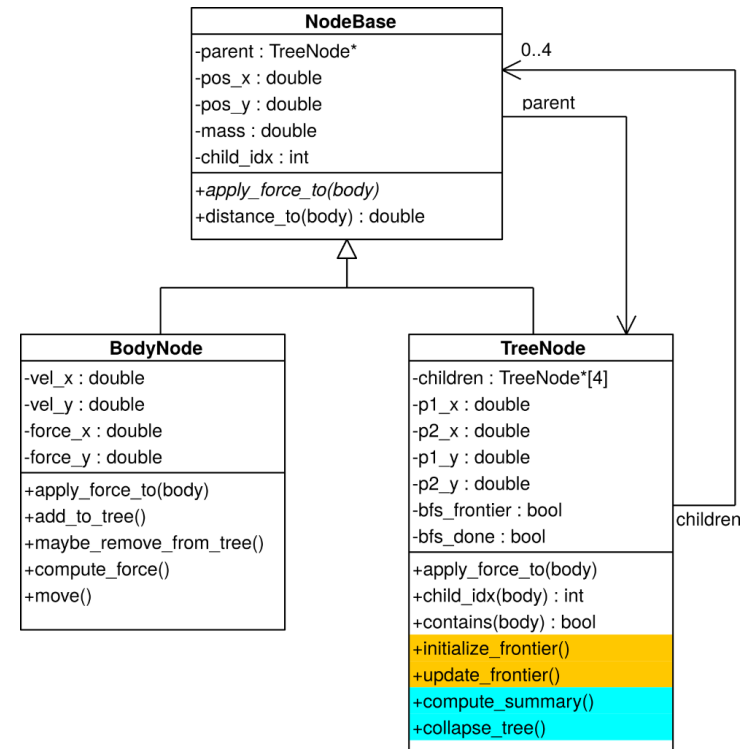



(h) water [6]:
Agent-based Sim.

Dynamic allocator enables parallel tree construction



Example 3: Data Structure



OOP helps programming complicated applications



Eg. Traffic simulation

- Road Segment
 - Regular / Junction Expressway / ...
- Vehicle
 - Regular / Bus / Emergency / ...



Dynamic memory allocation inside of GPU

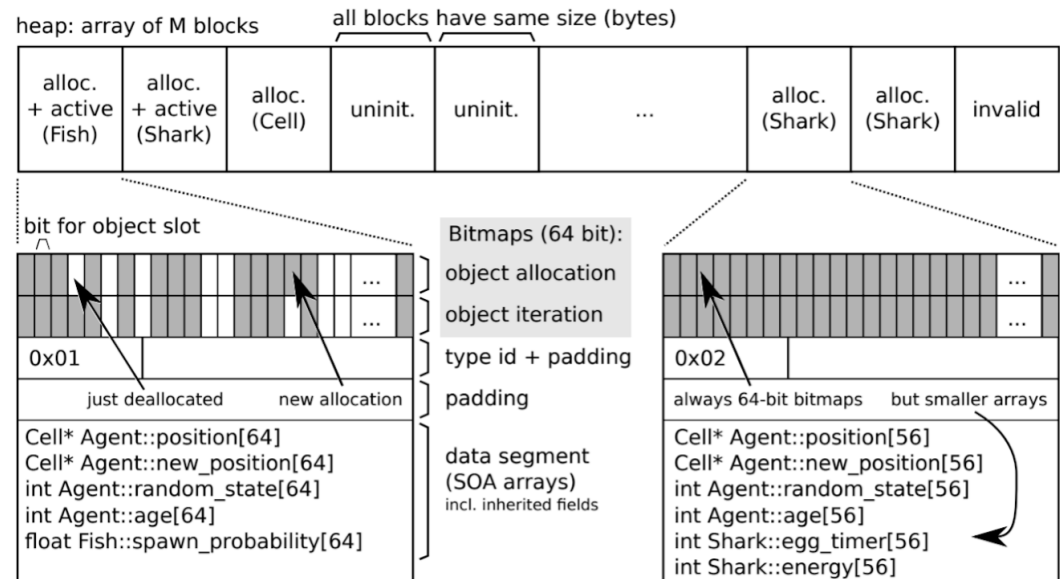



Challenges

- No global locks
- Keep allocated data dense

Approaches


- Block based
- Hierarchical bitmaps



Summary



- OOP helps programming complicated GPGPU applications



- A lot of “standard” mechanisms need to be reworked for GPGPU

