

SEBASTIAN BURCKHARDT



From Basel, Switzerland.

At Microsoft Research since 2007

General interest: programming models for distributed/parallel/concurrent systems.

- Particular interests:
 - Eventual Consistency (e.g. specifications, CRDTs)
 - Elastic services, actor systems (e.g. Orleans)
 - Serverless programming models

Foundations and Trends® in
Programming Languages
1:1-2

Principles of Eventual Consistency

Sebastian Burckhardt

BACKGROUND:

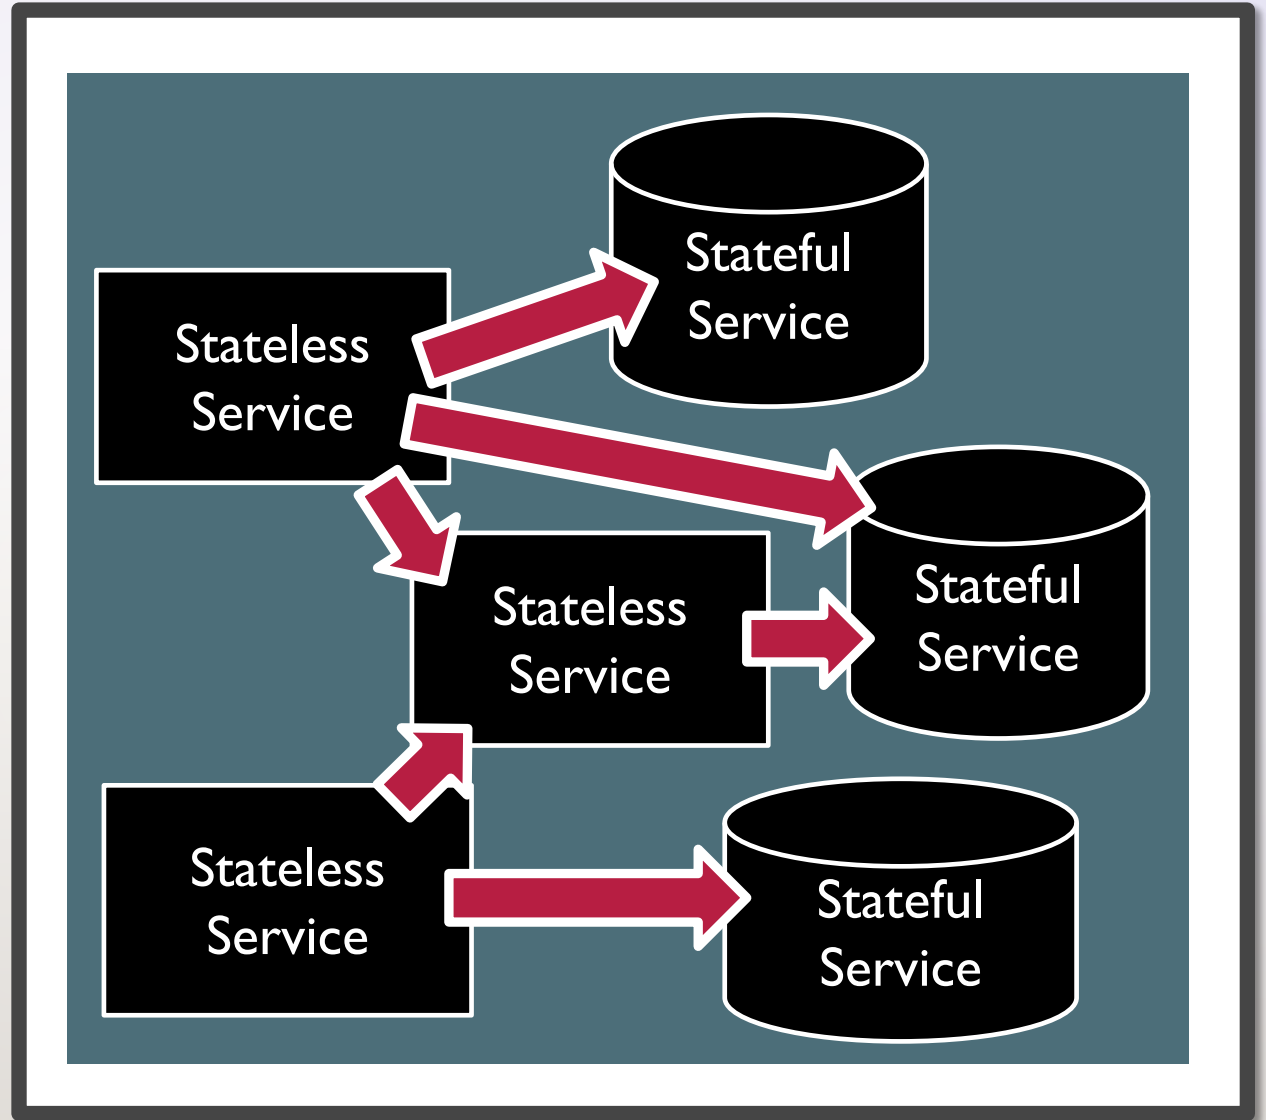
HOW TO BUILD SCALABLE, AVAILABLE SERVICES TODAY?

- From scalable, available components.
- Leave all the real hard stuff
(distributed protocols, failure detection, failure recovery,...)
for others to solve.

COMPOSED SERVICES

Services represent logical functionality, not physical machines.

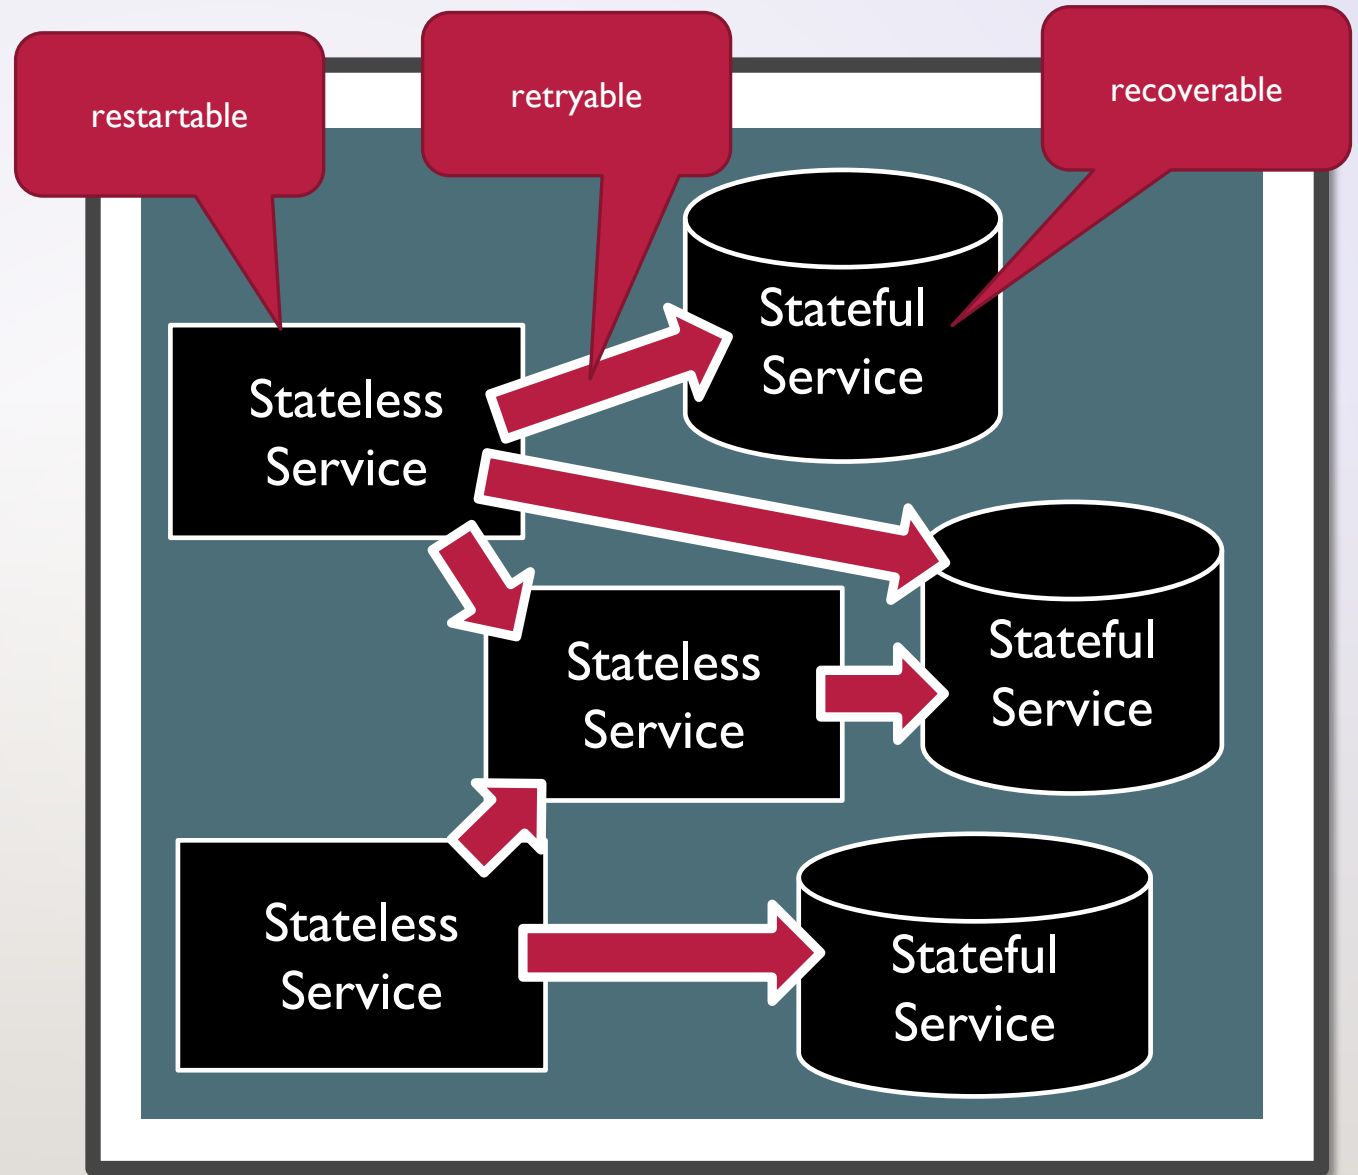
Often: multiple languages, REST APIs



COMPOSED SERVICES

All components are highly available.


Localized failure recovery via
idempotent or testable APIs.



INTERNAL PARTITIONING

services are **internally partitioned** so they can scale beyond the capacity of a single machine.

Partitions can be very fine-grained (e.g. key-value pair in a key-value store).




ELASTIC PACKING

Needed for
price/performance.

Much recent innovation in
this space.

General goals: make this

- a bit easier (K8s, Docker)
- fully automatic (serverless)



DEVELOPERS CHOOSE

Control

Productivity



Assembly

C,
C++

Java,
C#

Python

DEVELOPERS CHOOSE

Control

Productivity



Infrastructure
as a Service

Containers
as a service

Platform
as a service

Functions
as a Service

e.g.

AWS Lambda, Azure Functions

DEVELOPERS CHOOSE

SERVERLESS
(meaning: concept of a “server” is not visible in the application code, but managed within lower layer)

Control

Productivity



Infrastructure
as a Service

Containers
as a service

Platform
as a service

Functions
as a Service
e.g.
AWS Lambda, Azure Functions

TOP-GROWING CLOUD SERVICES 2019

| Place | Service | Growth | 2018 Use | 2019 Use |
|----------|------------------------|--------|----------|----------|
| #1 (tie) | Serverless | 50% | 24% | 36% |
| #1 (tie) | Stream Processing | 50% | 20% | 30% |
| #3 | Machine Learning | 44% | 18% | 26% |
| #4 | Container-as-a-Service | 42% | 26% | 37% |
| #5 | IoT | 40% | 15% | 21% |
| #6 | Data warehouse | 38% | 29% | 40% |
| #7 | Batch processing | 38% | 26% | 36% |

Source: Forbes, RightScale 2019 state of the cloud report

SERVERLESS FUNCTIONS

```
string helloworld(string name)
{
    return "Hello, " + name;
}
```

- Easy to deploy
- Elastic scale
- Load-based cost (e.g. pay per invocation)
- Free language choice, easy REST interface

```
> curl http://my-function-app.azure.com/helloworld?name=Shonan
Hello, Shonan
```

SERVERLESS FUNCTIONS ARE NOT “PURE”. THEY CAN CALL OTHER SERVICES.

Functions can **call** external services:

key-value stores, queues, blob storage,
pub-sub, databases, ...

= the “standard library” of cloud
programming!

```
async void delete_all()
{
    await cloudstorage.delete_file("*");
}
```

```
async void counter_increment()
{
    var current = await cloudstorage.read("counter");
    current = current + 1;
    await cloudstorage.write("counter");
}
```

Stateless Functions are great

but...

... WE SEE SOME PAIN POINTS AROUND STATE MANAGEMENT AND SYNCHRONIZATION.

- Synchronization

functions can interleave and race, synchronization via storage is challenging

- Partial execution

hosts can fail in the middle of a function, leaving behind inconsistent state

- Performance

Lots of calls to storage, lots of data movement => wastes time and CPU.

AZURE DURABLE FUNCTIONS

State & Synchronization for Serverless

2 NEW TYPES OF STATEFUL FUNCTIONS

Stateless
Functions



Orchestrations
(\approx Workflows)

provide

- durable **execution** state
- critical sections



Entities
(\approx Actors)

provide

- durable **application** state
- operation sequencing

ORCHESTRATIONS

= DURABLE EXECUTION STATE

- Reliably execute some combination of functions.
 - e.g. a simple sequence of functions, or multiple parallel function calls
- Eliminates the **partial execution** problem.

EXAMPLE I

- Upload file, then add to index

```
void upload_image(string name, byte[] data)
{
    await addtoblobstorage(name, data);

    await addtoindex(name);
}
```

```
void addtoblobstorage(string name, byte[] data)
{
    ...
}
```

```
void addtoindex(string name)
{
    ...
}
```


ORCHESTRATIONS

- Similar to workflows, but **straightforward async-await code**.
No XML or state machines.
- Thus, benefit from complete experience of the host language for control flow (sequential composition, parallel composition, all kinds of loops, exception handling, ...).

EXAMPLE 2

- Send many messages in parallel

```
void send_text_to_all_friends(string text)
{
    var friends = await getFriends();

    await Task.WhenAll(
        friends.Select(f => send(f, text)).ToList());
}
```

```
list<userid> getFriends()
{
    ...
}
```

```
send(userid destination, string text)
{
    ...
}
```

IMPLEMENTATION: HOW DOES IT WORK?

- Runtime uses record & replay
- Under failures, calls may be duplicated (but only the very last one)

```
void upload_image(string name, byte[] data)
{
    await addtoblobstorage(name, data);
    await addtoindex(name);
}
```

Log:

1. start function, input = (name, data)
2. call addtoblobstorage, input =(name, data)
3. response received, output = ()
4. call addtoindex, input = (name)
5. response received, output = ()
6. finish

IMPLEMENTATION: HOW DOES IT WORK?

- Runtime uses record & replay.
- Under failures, calls may be duplicated (but only the very last one)
- **Requires code to be deterministically replayable.**
Source of bugs (e.g. user must call context.UtcNow, not DateTime.UtcNow)
- **PL research could help:**
 - language could use type/effect system to track determinism / nondeterminism
 - language could provide serializable execution state, making replay unnecessary

```
void upload_image(string name, byte[] data)
{
    await addtoblobstorage(name, data);
    await addtoindex(name);
}
```

2 NEW TYPES OF STATEFUL FUNCTIONS

Stateless
Functions



Orchestrations
(\approx Workflows)

provide

- durable **execution** state
- critical sections



Entities
(\approx Actors)

provide

- durable **application** state
- operation sequencing

ENTITIES

= DURABLE APPLICATION STATE

- Entity = small piece of state identified by (name, key) string pair.
- Runtime delivers “operations” (messages) to entities via ordered async channels
- Runtime executes operations on entities, *one at a time*. Operations can
 - read and update state
 - send messages
 - perform external calls
 - return a value to caller (if the caller is an orchestration)
- **Durable**: All state (incl. messages) reliably kept in cloud storage

EXAMPLE ENTITY: BANK ACCOUNT

- each entity identified by a (name,key) pair, e.g. (“AccountEntity”, “32974-234093-00”)
- Serverless function defines how the entity handles operations, e.g. untyped interface in C#:

```
[FunctionName("AccountEntity")]
public static void Run([EntityTrigger] IDurableEntityContext context)
{
    switch(context.OperationName)
    {
        case "get":
            context.Return(context.GetState<int>());
            break;
        case "add":
            context.SetState(context.GetState<int>() + context.GetInput<int>());
            break;
    }
}
```

CALL VS. SIGNAL

- An actor can **signal** another actor
send message, fire and forget
- An orchestration can **call** an actor
and wait for ack/result
- But actors cannot call actors (to prevent deadlock)

FEATURE: EXPLICIT LOCKING

- Orchestrations can use **critical sections**
- very effective for preventing unwanted data races and interleavings
- Critical sections never fail spuriously and require no rollback or compensations

EXAMPLE: TRANSFER FUNDS

```
var fromAccount = new EntityId("AccountEntity", from);
var toAccount = new EntityId("AccountEntity", to);

using (await ctx.LockAsync(fromAccount, toAccount))
{
    var availablebalance = await ctx.CallEntityAsync<int>(fromAccount, "get");

    if (amount <= availablebalance)
    {
        await Task.WhenAll(
            ctx.CallEntityAsync(fromAccount, "add", -amount),
            ctx.CallEntityAsync(toAccount, "add", amount)
        );
    }
}
```



MESSAGE DIAGRAM

```
var fromAccount = new EntityId("AccountEntity", from);
var toAccount = new EntityId("AccountEntity", to);

using (await ctx.LockAsync(fromAccount, toAccount))
{
    var availablebalance = await ctx.CallEntityAsync<int>(

    if (amount <= availablebalance)
    {
        await Task.WhenAll(
            ctx.CallEntityAsync(fromAccount, "add", -amount),
            ctx.CallEntityAsync(toAccount, "add", amount)
        );
    }
}
```

orchestration



DEADLOCK PREVENTION

We enforce some simple rules to prevent deadlocks:

- Runtime acquires locks in order (fixed global total order).
- Critical sections cannot be nested.
- Within a critical section:
 - can call only entities that were locked.
 - can signal only entities that were not locked.

STATUS

- Azure Durable Functions have been out for about a year now.
(thanks to Chris Gillum & Durable Functions team)
- Entities (& critical sections) are a new feature I implemented over the last couple months, building on research done last year w/ intern Christopher Meiklejohn, **now in public preview since last month.**
- No paper yet.