



# Virtual Machine Support for Debugging Concurrency Bugs

No Distribution!

Deterministic Replay, Asynchronous Snapshots,  
and Bug Detection for Event Loop Systems

**Stefan Marr**

Shonan, May 2019

Ongoing work with Dominik Aumayr, Carmen Torres Lopez,  
Elisa Gonzalez Boix, Hanspeter Mössenböck

University of  
**Kent**

# Ongoing work



involving



Carmen  
Torres Lopez



Elisa  
Gonzalez Boix



Dominik  
Aumayr



Hanspeter  
Mössenböck



Based on **GraalVM**™

# Concurrency Bugs are Common in Event Loop Systems



53 projects, 57 issues

2 studies

12 projects, 1000 potential issues



12 projects

1 study

53 concurrency issues



Websites in top 500

≈1-10 concurrency issues



8-27 apps

≈2-20 concurrency issues



6 projects

35 known event races



# Challenges



## bug reproduction

e.g. fails 1 in 10 times



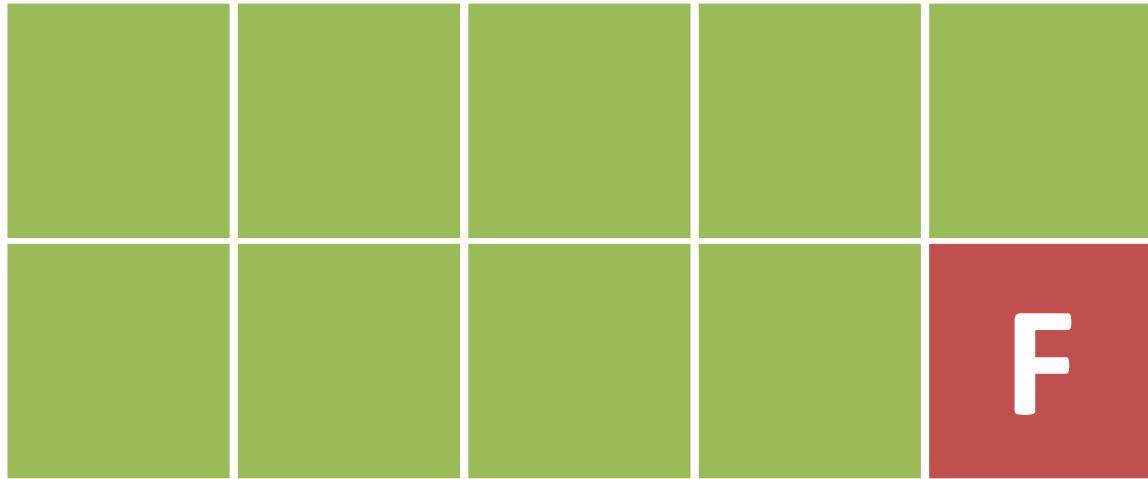
## asynchronous snapshots

saving the world, without stopping it



## bug mitigation

fails 1 in 10 times, can we avert failure?

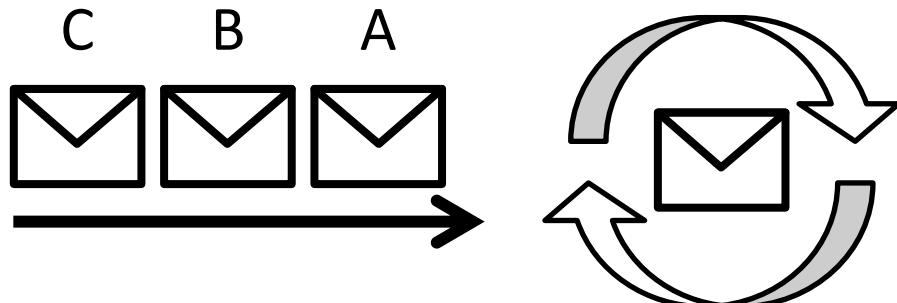


Reproduces only 1 in 10? How can I fix such a bug???

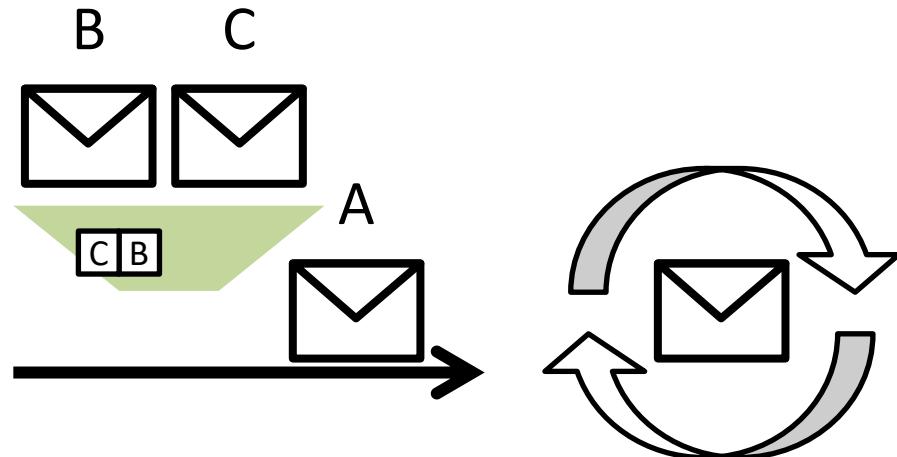
# **NON-DETERMINISM MAKES FOR UNHAPPY DEBUGGERS**

# One Solution

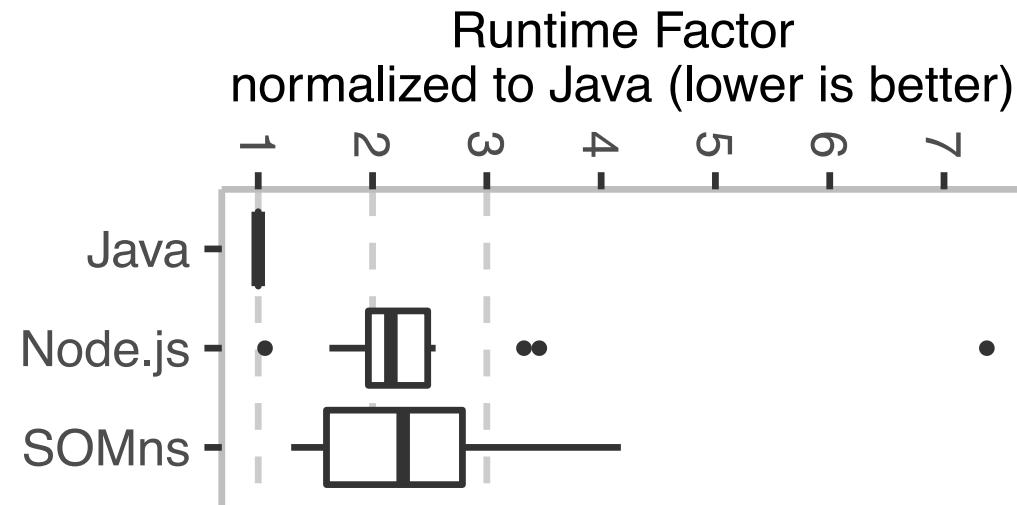
- Record event order



- Replay reorder to fit

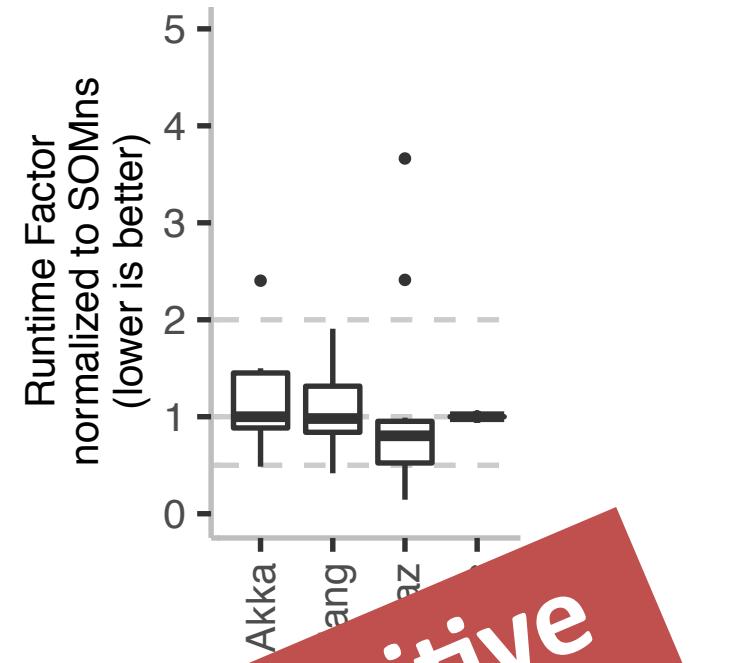


# Performance: Baselines



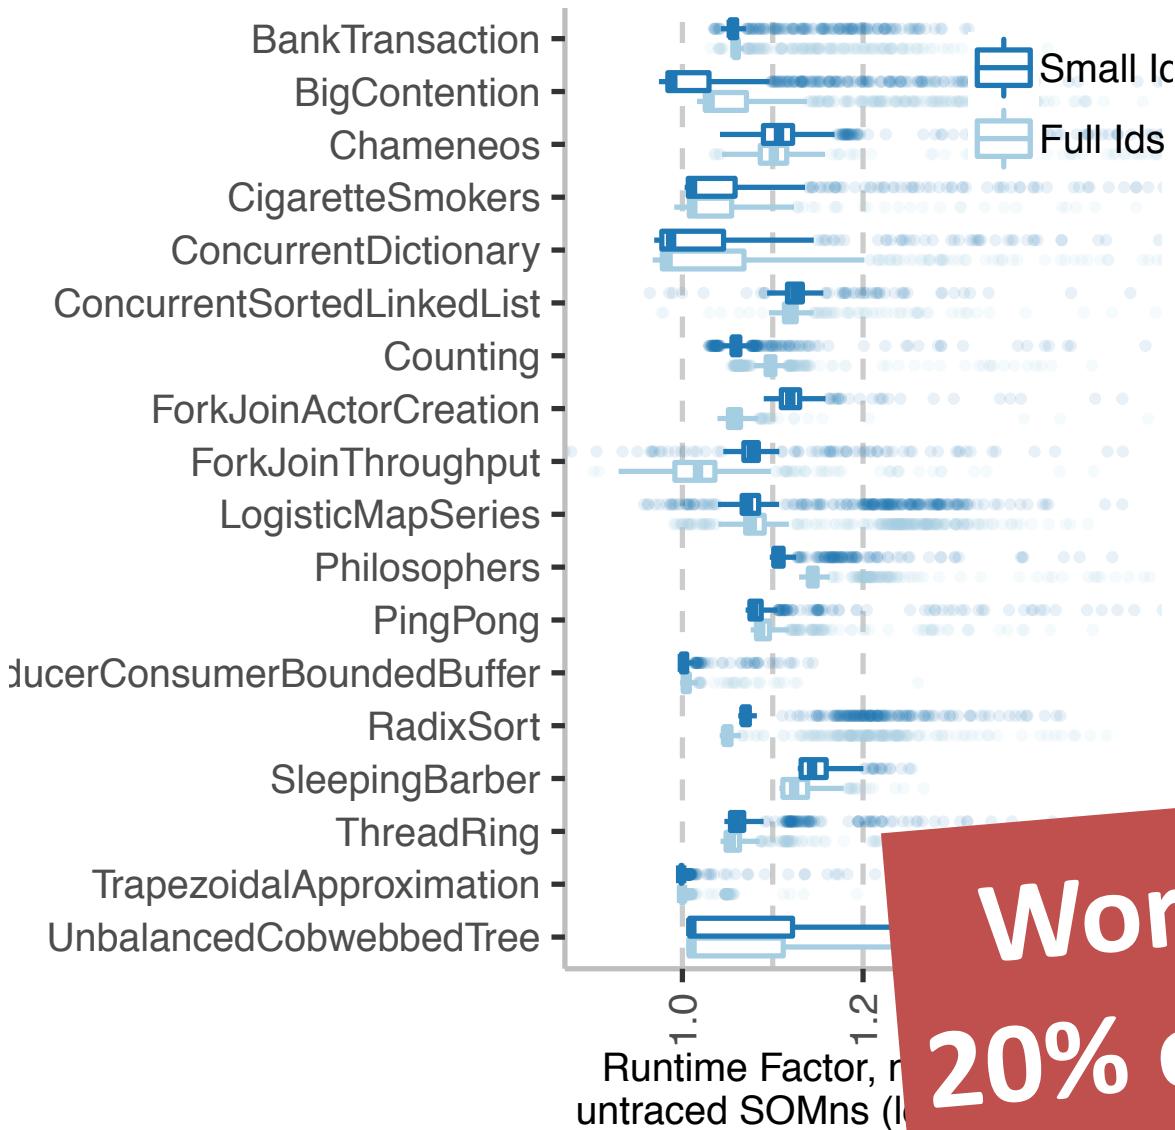
*On level of  
Dynamic  
Languages!*

Cross-Language Comparison  
[smarr/are-we-faster-than-jvm](https://github.com/smarr/are-we-faster-than-jvm)  
Benchmark Suite  
[shamsimam/savina](https://github.com/shamsimam/savina)



*Competitive  
with JVM  
frameworks!*

# Tracing Performance: Savina

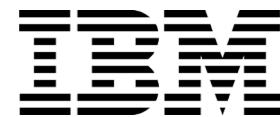


Worst case:  
20% overhead



# AcmeAir: Microservices Evaluator

- Airline booking system
- Designed as benchmark

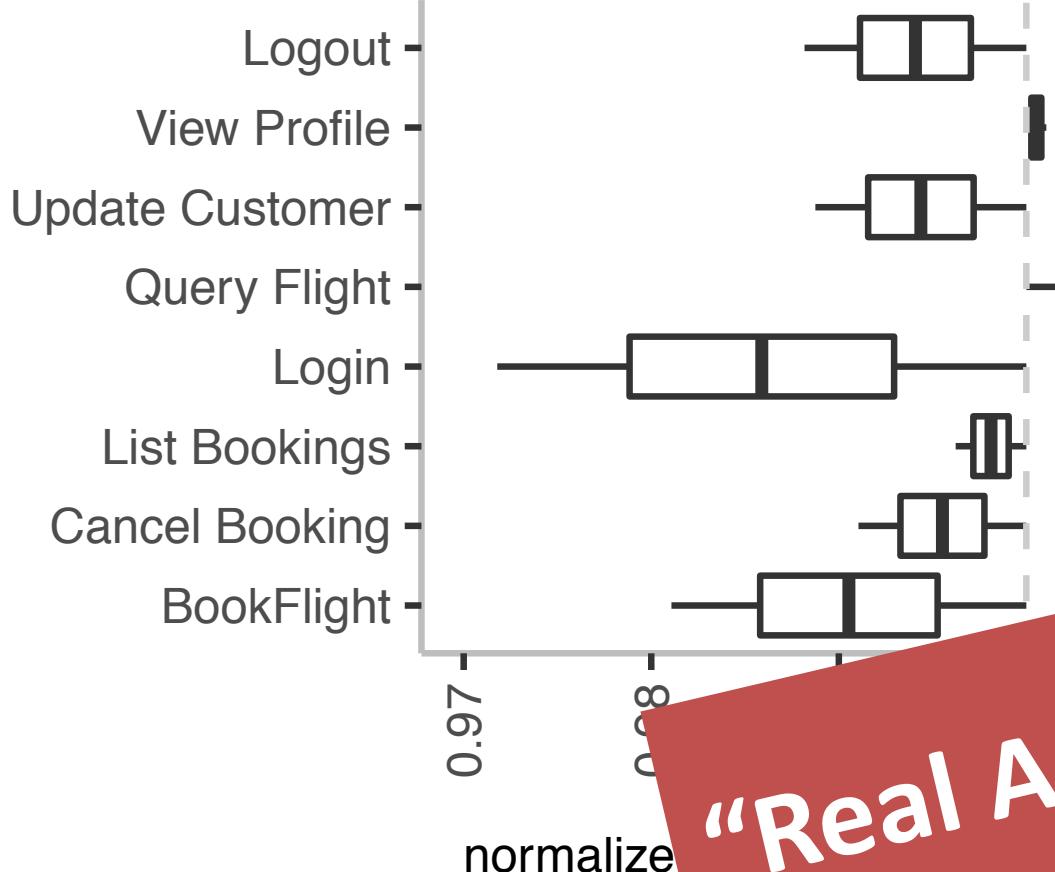


<https://github.com/acmeair>





# Tracing Performance: Acme Air



“Real Application”  
impact: very low



Snapshotting Actor Systems without Stopping Them

**LONG AND HUGE TRACES MAKE  
REPLAY IMPRACTICAL**

# Deterministic Replay and Snapshots

## Microservices in the Cloud

Service Overview

As of Nov 16, 2017 9:52:31 PM UTC

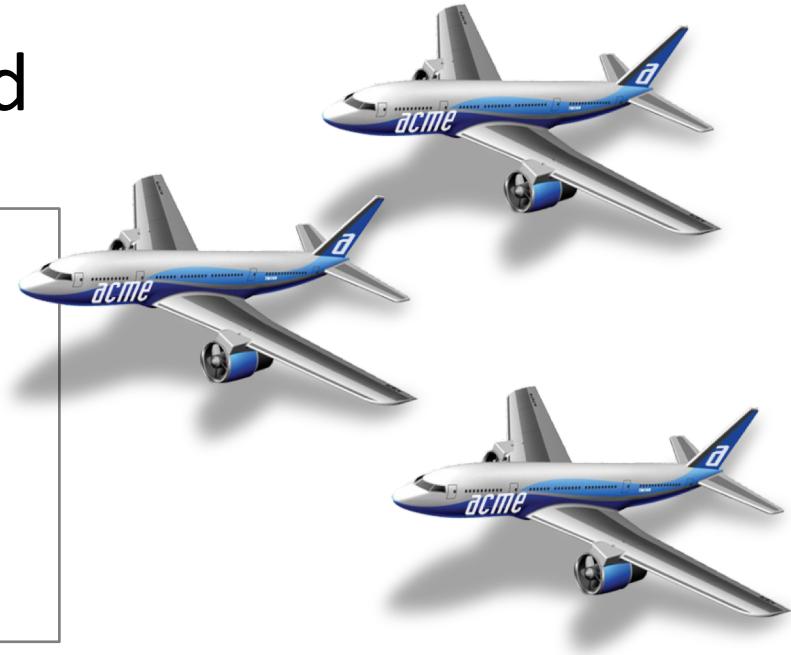
1	2	30 GB	77 GB
Nodes	OCPUs	Memory	Storage

Status: Creating service ...  
Backup Destination: BOTH  
Enable B2B adapter for EDI: true  
Open Sample Application: [https://\[REDACTED\]/sample-app](https://[REDACTED]/sample-app)

Version: 12.2.1.2.0  
Cloud Storage Container: [REDACTED]  
JDK: 1.8.0\_144  
Service Type: SOA with SB & B2B

Resources

Host Name: myoraclesoaservice-wls-1	OCPUs: 2
Public IP: [REDACTED]	Memory: 30 GB
Instance: Runs MyOracle_server_1	Storage: 77 GB



Service runs for days or weeks  
Can't have replay from start → Need Snapshots

# Deterministic Replay and Snapshots

## Microservices in the Cloud



Service Overview As of Nov 16, 2017 9:52:31 PM UTC

1 Nodes	2 OCPUs	30 GB Memory	77 GB Storage
---------	---------	--------------	---------------

Status: Creating service ... Version: 12.2.1.2.0  
Backup Destination: BOTH Cloud Storage Container:   
Enable B2B adapter for EDI: true JDK: 1.8.0\_144  
Open Sample Application: [/sample-a">https:///sample-a](https://<img alt=)

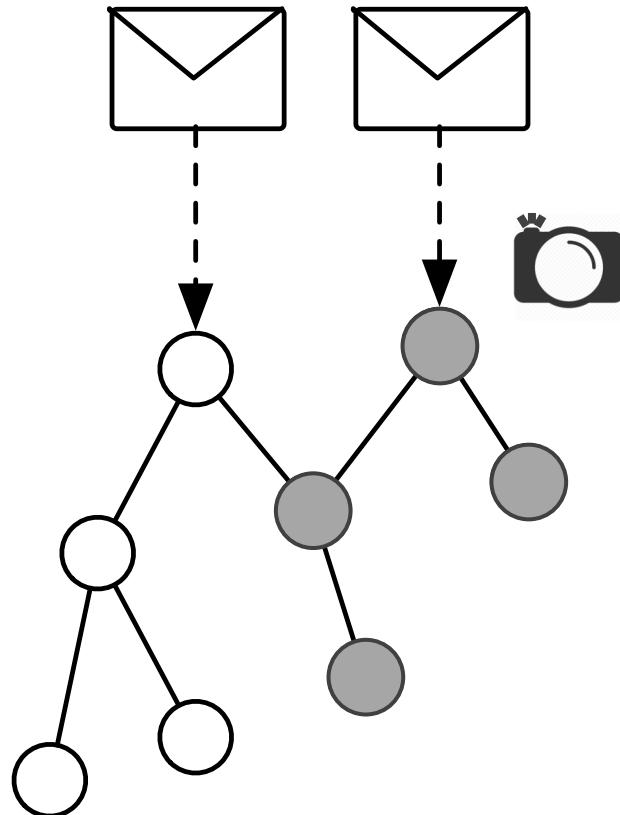
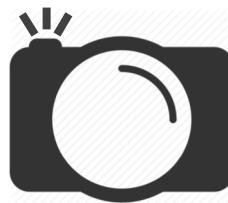
Resources

Host Name:	myoraclesoaservice-wls-
Public IP:	
Instance:	Runs MyOracle_server_1

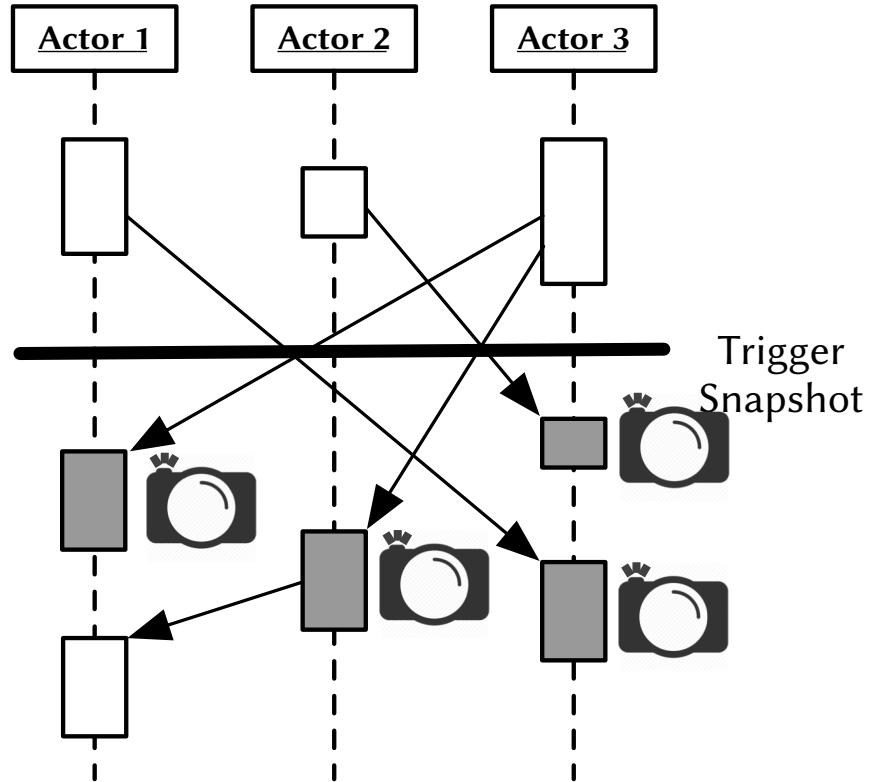
**Download Snapshot.  
Debug Locally!**



# Asynchronous and Partial Heap Snapshots



snapshot only objects  
reachable from a message



serialize partial heap  
before message execution



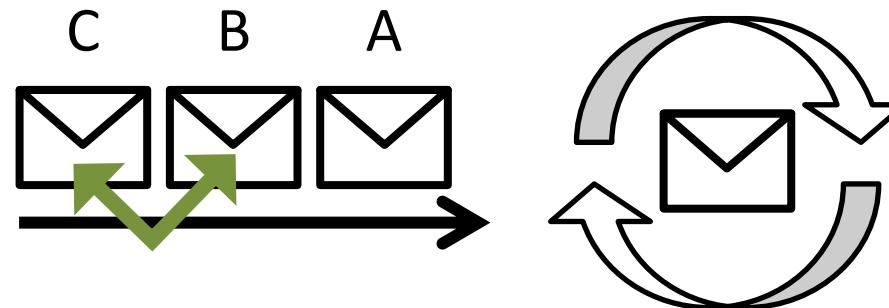
Where to go from here?

# **THE FUTURE?**



# Bug Mitigation?

## Detect Event Races At Run Time



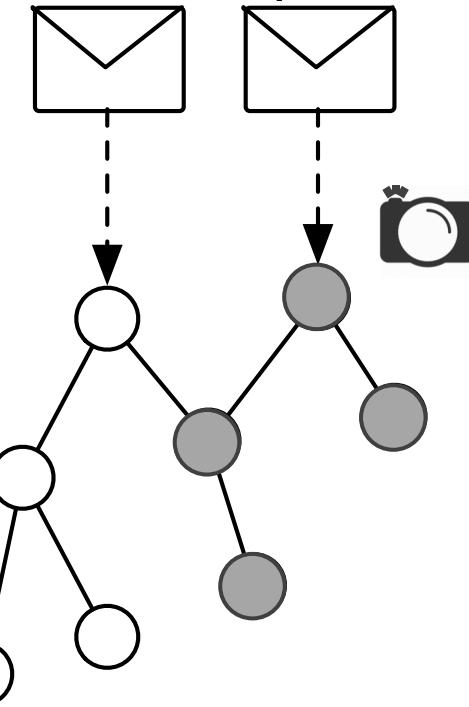
Order A -> B -> C problematic?

Let's swap them!

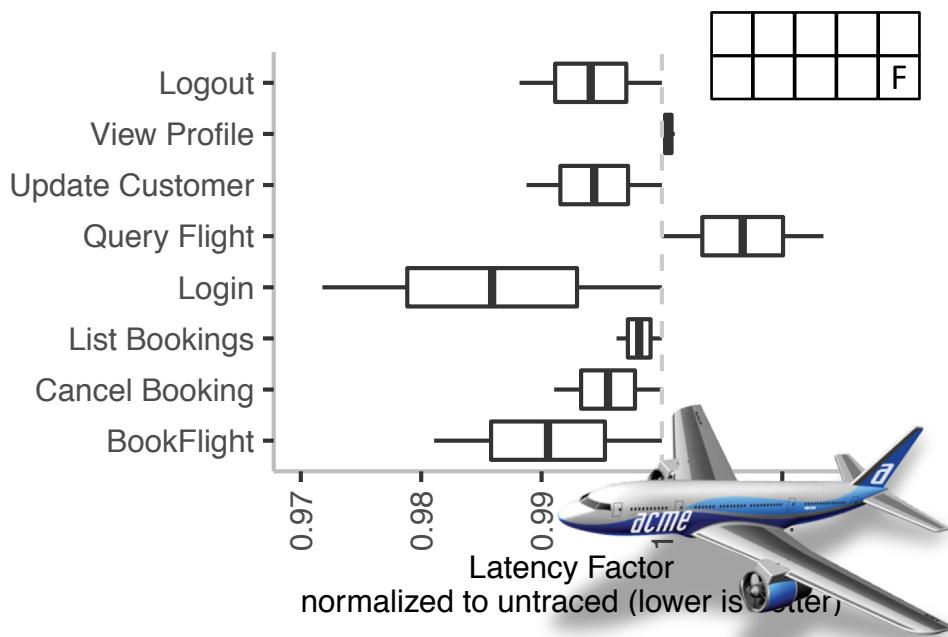
# Summary



## Asynchronous Partial Snapshots



## Deterministic Replay with low overhead



Future: Race Detection and Bug Mitigation

